

CHENGSHUAI ZHAO

Ph.D. Student in Computer Science · Arizona State University · Tempe, AZ, USA

☎ +1 (949) 227-7633 | ✉ czhao93@asu.edu | 🌐 chengshuaizhao0.github.io | 🔗 LinkedIn

EDUCATION

Arizona State University <i>Ph.D. in Computer Science, Advisor: Huan Liu</i>	Jan 2024 – Jan 2029 (expected) Tempe, AZ, USA
University of California, Irvine <i>B.S. in Data Science</i>	Sep 2022 – Jun 2023 Irvine, CA, USA
Huazhong Agricultural University <i>B.S. in Data Science · B.S. in Life Science</i>	Sep 2018 – Jun 2023 Wuhan, China

RESEARCH EXPERIENCE

Data Mining and Machine Learning Lab (DMML), Arizona State University <i>Graduate Research Assistant, advised by Prof. Huan Liu</i>	Jan 2024 – Present Tempe, AZ, USA
---	--------------------------------------

- **LLM post-training & reasoning:** reasoning distillation via explanatory data synthesis and RL with structured rewards [4]; an RL-optimized knowledge graph retriever for reasoning-augmented generation [12]; a controlled and synthetic environment for pre-training LLMs from scratch and probing reasoning [7].
- **LLM/MLLM system security, privacy & safety:** bilevel-optimized knowledge-poisoning attacks on RAG systems [19] and covert malicious-agent injection into MLLM societies [18]; adversarial data synthesis that protects data from unauthorized fine-tuning of LLMs [28] and VLMs [29]; RL-driven defenses that shield proprietary LLMs from unauthorized API-based knowledge distillation [30].
- **RAG & LLM agent applications:** knowledge-graph-grounded response verification in RAG for domain-specific question answering [13], [14]; a multi-agent RL framework in Graph-RAG for multi-hop reasoning [12]; LLM multi-agent frameworks with human-in-the-loop collaboration for large-scale content analysis [15], [16].
- **LLM evaluation:** LLM-as-a-judge [10] and the detectability of LLM-generated judgments [11]; detection and evaluation of LLM-generated content [17], [31]; evaluation of CoT reasoning under distribution shift [7].
- **Causality, interpretability, efficiency & recommendation:** causality-guided representation learning for robust generalization across domains, styles, and modalities [3], [8], [27]; concept bottleneck models for interpretable and trustworthy open-ended grading [9], [1]; low-rank node-level graph diffusion for few-shot node classification [6]; diffusion language models for diverse and efficient bundle recommendation [26].

Center for AI Research and Innovation (CAIRI), Westlake University <i>Research Intern, advised by Prof. Stan Z. Li</i>	Jul 2022 – Sep 2022 Hangzhou, China
--	--

- **Molecular foundation models:** self-supervised pre-training of GNNs for molecules with a context-aware atom-level tokenizer and masked atom modeling with contrastive learning [20].

BioMedical Big Data Mining Lab (BBDM), Huazhong Agricultural University <i>Undergraduate Researcher, advised by Prof. Wen Zhang</i>	Oct 2019 – Sep 2022 Wuhan, China
---	-------------------------------------

- **Graph representation learning & link prediction:** contrastive self-supervised GNNs, equivariant positional-encoding GNNs, and graph-embedding ensembles for predicting molecular interactions [21], [23], [24].

CONFERENCE, JOURNAL & WORKSHOP PAPERS

- [1] K. S. Chaudhary, **C. Zhao**, F. Zhang, G. Agrawal, Y. Deng, and H. Liu, "EssayCBM: Rubric-Aligned Concept Bottleneck Models for Transparent Essay Grading," **ECML-PKDD 2026**.
- [2] Y. Li, **C. Zhao**, G. Agrawal, Y. Deng, J. Park, Y.-C. Chen, and H. Liu, "Epistemic Cognition and Uncertainty Navigation with a Domain-Specific AI Chatbot in STEM Education," **LERN 2026**.
- [3] P. Ma, **C. Zhao**, B. Jiang, S. Vishnubhatla, U. Jeong, A. Beigi, A. Raglin, and H. Liu, "CAMO: Causality-Guided Adversarial Multimodal DOMain Generalization for Crisis Classification," **PAKDD 2026 (Oral)**.
- [4] Z. Tan, **C. Zhao**, S. Wang, J. Li, T. Chen, and H. Liu, "Probing to Refine: Reinforcement Distillation of LLM Reasoners via Explanatory Inversion," **ICLR 2026**.
- [5] S. Wang, Z. Tan, Z. Chen, D. Li, Y. Zhu, B. Jiang, Y. He, **C. Zhao**, Z. Lei, P. Sheth, L. Li, L. P.-Y. Ting, J. Li, and H. Liu, "Large Language Models for Data Science: A Survey," **SIGKDD Explorations, 2026**.
- [6] Y. Wang, **C. Zhao**, D. Sun, H. Liu, and Y. Yang, "Low-Rank Few-Shot Node Classification by Node-Level Graph Diffusion," **ICLR 2026**.
- [7] **C. Zhao**, Z. Tan, P. Ma, D. Li, B. Jiang, Y. Wang, Y. Yang, and H. Liu, "Is Chain-of-Thought Reasoning of LLMs a Mirage? A Data Distribution Lens," **ACL 2026**.
- [8] **C. Zhao**, S. Wan, P. Sheth, K. Patwa, K. S. Candan, and H. Liu, "Causality Guided Representation Learning for Cross-Style Hate Speech Detection," **WWW 2026 (Oral)**.
- [9] **C. Zhao**, F. Zhang, K. S. Chaudhary, Y. Li, L. P.-Y. Ting, Y.-C. Chen, and H. Liu, "REC-CBM: Rubric-Aware Error-Correction Concept Bottleneck Models for Trustworthy Open-Ended Grading," **SDM 2026**.
- [10] D. Li, B. Jiang, L. Huang, A. Beigi, **C. Zhao**, Z. Tan, A. Bhattacharjee, Y. Jiang, C. Chen, T. Wu, K. Shu, L. Cheng, and H. Liu, "From Generation to Judgment: Opportunities and Challenges of LLM-as-a-judge," **EMNLP 2025 (Oral)**.
- [11] D. Li, Z. Tan, **C. Zhao**, B. Jiang, B. Huang, P. Ma, A. Alnaibari, K. Shu, and H. Liu, "Who's Your Judge? On the Detectability of LLM-Generated Judgments," **Lock-LLM @ NeurIPS 2025**.
- [12] L. P.-Y. Ting*, **C. Zhao***, Y.-H. Zeng, Y. J. Lim, K.-T. Chuang, and H. Liu, "Leaps Beyond the Seen: Reinforced Reasoning Augmented Generation for Clinical Notes," **GenAI4Health @ NeurIPS 2025**.
- [13] **C. Zhao**, G. Agrawal, F. Zhang, T. Kumarage, Z. Tan, Y. Deng, Y. Chen, and H. Liu, "Ontology-Aware RAG for Improved Question-Answering in Cybersecurity Education," **IEEE BigData 2025**.
- [14] **C. Zhao**, R. D. Maria, T. Kumarage, K. S. Chaudhary, G. Agrawal, Y. Li, J. Park, Y. Chen, Y. Deng, and H. Liu, "CyberBOT: Ontology-Grounded Retrieval Augmented Generation for Reliable Cybersecurity Education," **CIKM 2025 (Demo Track)**. 🏆 **Best Paper Award**.
- [15] **C. Zhao**, Z. Tan, C. Wong, X. Zhao, T. Chen, and H. Liu, "SCALE: Towards Collaborative Content Analysis in Social Science with Large Language Model Agents and Human Intervention," **ACL 2025 (Oral)**.
- [16] X. Zhao, **C. Zhao**, M. Mengesteab, C.-W. Wong, M. Zhan, Z. Tan, and T. Chen, "Automating Content Analysis With Multiple LLM Agents: Impacts of Agent Attributes and Human-AI Collaboration," **Social Science Computer Review, 2025**.
- [17] B. Jiang*, **C. Zhao***, Z. Tan, and H. Liu, "Catching Chameleons: Detecting Evolving Disinformation Generated using Large Language Models," **IEEE CogMI 2024**. 🏆 **Best Student Paper Award**.
- [18] Z. Tan*, **C. Zhao***, R. Moraffah, Y. Li, Y. Kong, T. Chen, and H. Liu, "The Wolf Within: Covert Injection of Malice into MLLM Societies via an MLLM Operative," **ReGenAI @ CVPR 2024**.

- [19] Z. Tan*, **C. Zhao***, R. Moraffah, Y. Li, S. Wang, J. Li, T. Chen, and H. Liu, “Glue pizza and eat rocks - Exploiting Vulnerabilities in Retrieval-Augmented Generative Models,” **EMNLP 2024**.
- [20] J. Xia*, **C. Zhao***, B. Hu, Z. Gao, C. Tan, Y. Liu, S. Li, and S. Z. Li, “Mole-BERT: Rethinking Pre-training Graph Neural Networks for Molecules,” **ICLR 2023**.
- [21] **C. Zhao**, H. Wang, W. Qi, and S. Liu, “Toward drug-miRNA resistance association prediction by positional encoding graph neural network and multi-channel neural network,” **Methods, 2022**.
- [22] X. Xu, S. Liu, Z. Yang, X. Zhao, Y. Deng, G. Zhang, J. Pang, **C. Zhao**, and W. Zhang, “A systematic review of computational methods for predicting long noncoding RNAs,” **Briefings in Functional Genomics, 2021**.
- [23] **C. Zhao**, S. Liu, F. Huang, S. Liu, and W. Zhang, “CSGNN: Contrastive Self-Supervised Graph Neural Network for Molecular Interaction Prediction,” **IJCAI 2021**.
- [24] **C. Zhao**, Y. Qiu, S. Zhou, S. Liu, W. Zhang, and Y. Niu, “Graph embedding ensemble methods based on the heterogeneous network for lncRNA-miRNA interaction prediction,” **BMC Genomics, 2020**.

PREPRINTS & PAPERS UNDER REVIEW

- [25] Z. Liu, **C. Zhao**, and H. Liu, “MUGEN: Generating Unlearnable Graph Examples for Multiple Learning Tasks,” *Under review at WSDM 2027*.
- [26] P. Ma, B. Jiang, **C. Zhao**, D. Li, Z. Tan, and H. Liu, “Diversity-Preserving and Efficient Diffusion Language Model Recommendation for Bundle Construction,” *Under review at KDD 2027*.
- [27] P. Ma, **C. Zhao**, U. Byahut, B. Jiang, S. Vishnubhatla, U. Jeong, A. Raglin, A. Harrison, and H. Liu, “CaDiF: Few-Shot Multimodal Learning via Causality-Based Disentanglement and Cross-Modal Flattening,” *Under review at WSDM 2027*.
- [28] **C. Zhao**, P. Ma, D. Li, B. Jiang, Z. Yu, Z. Tan, and H. Liu, “TextCloak: Thwarting Unauthorized LLM Exploitation via RL-Driven Unlearnable Text,” *Under review at AAAI 2027*.
- [29] **C. Zhao**, Z. Tan, D. Li, Z. Yu, and H. Liu, “To See is Not to Learn: Protecting Multimodal Data from Unauthorized Fine-Tuning of Large Vision-Language Model,” *Under review at NDSS 2027*.
- [30] **C. Zhao**, F. Zhang, D. Li, B. Jiang, P. Ma, Z. Yu, Z. Tan, and H. Liu, “Rethinking Anti-distillation of Large Language Models: An Unlearnable Examples Perspective,” *Under review at ICLR 2027*.
- [31] B. Jiang, D. Li, Z. Tan, **C. Zhao**, and H. Liu, “Are Today’s LLMs Ready to Explain Well-Being Concepts?” *Under review at AAAI 2027*.

HONORS & AWARDS

AWARDS

- **SIAM Student Travel Award**, SDM 2026 Sep 2026
- **Gold Reviewer Award**, ICML 2026 May 2026
- **SCAI Conference Award**, Arizona State University May 2026
- **Best Paper Award**, ACM CIKM 2025, Demo Papers Track Nov 2025
- **Best Student Paper Award**, IEEE CogMI 2024 Oct 2024
- **Innovation and Entrepreneurship Award**, Huazhong Agricultural University Nov 2021

SCHOLARSHIPS

- **Outstanding Study Scholarship** (top 0.1%) Jul 2021, Jul 2022
- **Advanced Study Scholarship** (top 3%) Apr 2019, Jul 2019, Jul 2020
- **Excellent Study Scholarship** (top 10%) Nov 2020, Nov 2021, Nov 2022

- **Merit Student Scholarship** (top 20%) Nov 2020

INVITED TALKS

- *Invited Talk*, NICE AI Talk, Virtual Dec 2025
- *Invited Talk*, CCSB Biweekly Seminar Series, Prairie View A&M University, Virtual Sep 2025
- *Guest Lecture*, CSE 472: Social Media Mining, Arizona State University, Tempe, AZ Aug 2025
- *Invited Talk*, Yochan Lab, Arizona State University, Tempe, AZ Aug 2025

MEDIA COVERAGE

- *The New Yorker*: What if A.I. Doesn't Get Much Better Than This? Aug 2025
- *Ars Technica*: LLMs' "Simulated Reasoning" Abilities Are a "Brittle Mirage," Researchers Find Aug 2025
- *VentureBeat*: LLMs Generate 'Fluent Nonsense' When Reasoning Outside Their Training Zone Aug 2025
- *SF Examiner*: New Study Calls AI Reasoning a 'Brittle Mirage' Aug 2025
- *The Decoder*: Yet Another Study Doubts That LLM Reasoning Shows True Logic over Pattern Imitation Aug 2025

PROFESSIONAL SERVICE

MEMBERSHIPS

- IEEE · ACM · AAAI · SIAM · AAAS

PROGRAM COMMITTEE MEMBER / REVIEWER

- **Data Mining & Information Retrieval**: SIGKDD, WWW, SIGIR, WSDM, CIKM, ICDM. 2024–2026
- **Machine Learning & AI**: ICML, NeurIPS, ICLR, AAAI, IJCAI. 2025–2026
- **Natural Language Processing**: ACL, EMNLP, NAACL, EACL. 2024–2025

WORKSHOP ORGANIZATION

- *Organizer & Speaker*, Workshop on Cybersecurity, AI, and Education, NSF SFS PI Meeting Feb 2026

TECHNICAL SKILLS

- **Research Skills**: LLM post-training (SFT, RL, knowledge distillation), LLM reasoning, LLM/MLLM system security, privacy & safety, LLM agents, RAG, LLM evaluation, explainable AI & model interpretability, graph representation learning (link prediction, node classification, graph generation), causal representation learning
- **Programming Languages**: C, C++, Java, Python, R, Go, Scala, SQL, Verilog HDL
- **Frameworks**: PyTorch, HF Transformers, DeepSpeed, vLLM, TensorFlow, Scikit-learn, Hadoop, Spark
- **Databases**: Oracle, MySQL, SQLite, MongoDB
- **Tools**: Git, Docker, Conda, Overleaf
- **Operating Systems**: Ubuntu, CentOS, macOS, MS Windows